How estimates change with chains per condition

Bold: reviewed model means · Faint: individual reviewed runs · At 50 chains: same means as the statistics table

H1 Outstanding-order sensitivity

Predicted: below zero

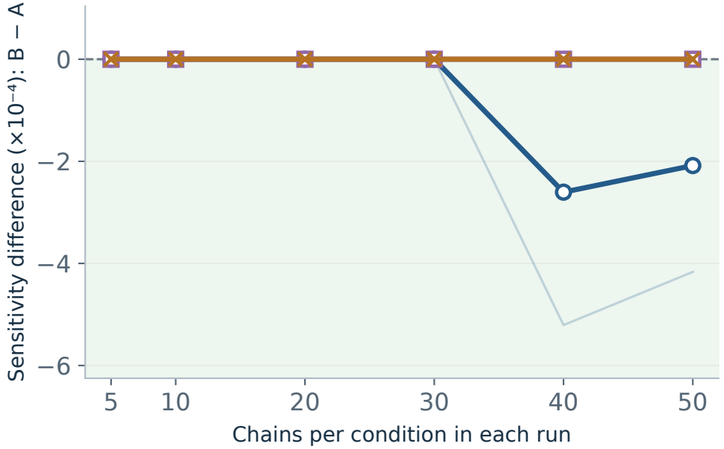

H2 Total inventory + backlog cost

Predicted: above zero

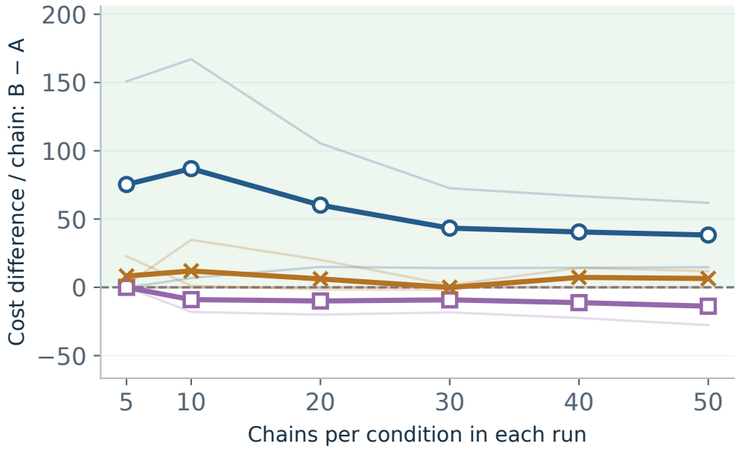

H3 Outstanding-order sensitivity

Predicted: above zero

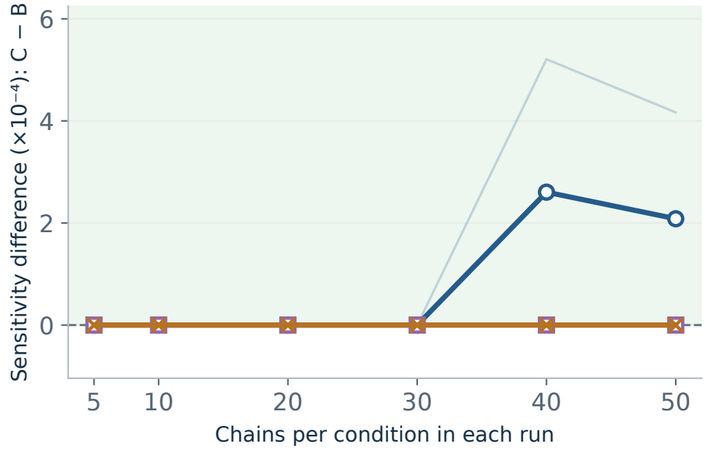

H4a Individual order variation

Predicted: below zero

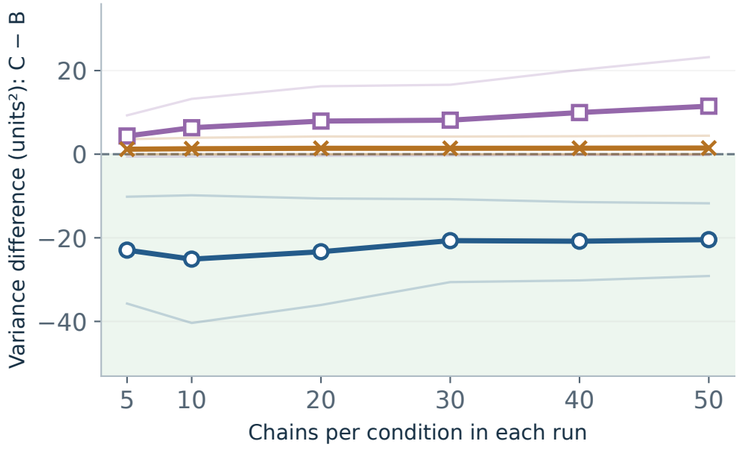

H4b Bullwhip ratio

Predicted: below zero

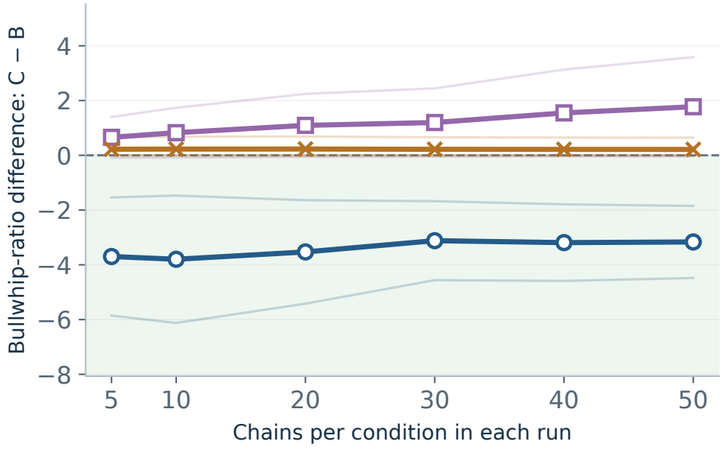

Same sample as the statistics

- GPT-6 Astra (n = 2 contexts)
- GPT-6.1 Sol (n = 2 contexts)
- GPT-6 Luna (n = 3 contexts)

How to read the axes

X: first k chains in each run
Independent n stays at 2 / 2 / 3
Green: predicted effect direction
H1 / H3: magnified $\times 10^{-4}$ scale

A = control + revision B = control + one decision C = causal education + one decision

Reviewed runs only. H1/H3 differences are tiny in absolute terms. Chain prefixes describe subset stability; no extra independent rounds are added.