Observed model differences across repeated runs

One point = one completed context, with 50 chains per condition and 24 weeks.

H2 Cost difference

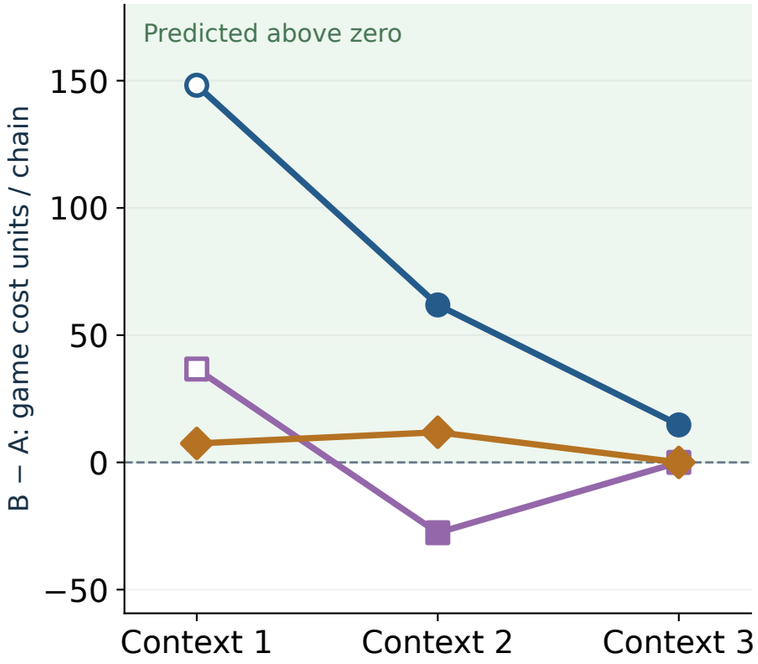

H4a Order-variance difference

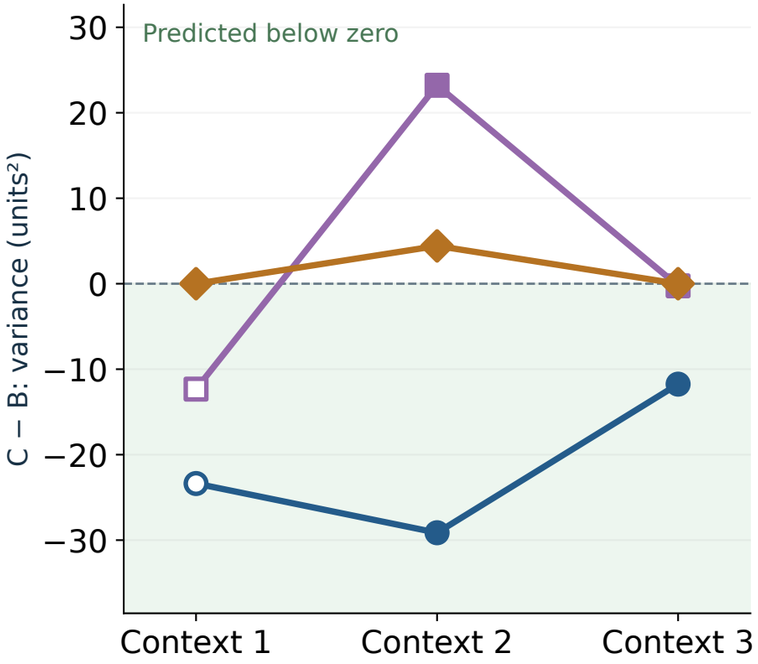

H4b Bullwhip difference

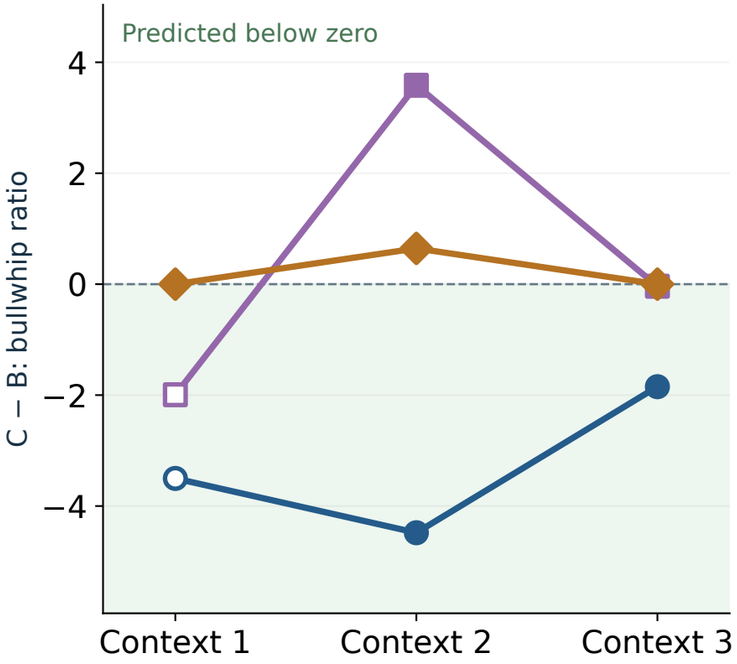

GPT-6 Astra GPT-6.1 Sol GPT-6 Luna

Filled markers: procedure reviewed. Open markers: provenance review incomplete. Excluded runs are omitted.

A = control + revision B = control + one decision C = causal education + one decision

Descriptive AI simulation. Reviewed-context effect sizes and exact p-values are provided in the H1–H4 summary slide.