H1–H4: observed effects and p-values

Reviewed contexts only. Δ = mean difference; p = raw two-sided p; Holm = adjustment across 20 tests.

Outcome	Predicted	GPT-6 Astra n = 1 contexts	GPT-6.1 Sol n = 2 contexts	GPT-6 Luna n = 3 contexts
H1 Constraint → pipeline sensitivity	Below 0	Δ +0 p — | Holm —	Δ +0 p 1 | Holm 1	Δ +0 p 1 | Holm 1
H2 Constraint → total cost	Above 0	Δ +61.9 p — | Holm —	Δ -13.8 p 1 | Holm 1	Δ +6.43 p 0.5 | Holm 1
H3 Education → pipeline sensitivity	Above 0	Δ +0 p — | Holm —	Δ +0 p 1 | Holm 1	Δ +0 p 1 | Holm 1
H4a Education → order variance	Below 0	Δ -29.1 p — | Holm —	Δ +11.5 p 1 | Holm 1	Δ +1.47 p 1 | Holm 1
H4b Education → bullwhip	Below 0	Δ -4.48 p — | Holm —	Δ +1.77 p 1 | Holm 1	Δ +0.215 p 1 | Holm 1

H1 / H3: no observed treatment differences in the reviewed pipeline-sensitivity measurements.

H2 / H4: observed signs differ across models and repeated contexts.

A = control + revision B = control + one decision C = causal education + one decision

n = reviewed independent contexts. — = p unavailable with one context. H1/H3 use pipeline sensitivity as a proxy.