Experiment 14: current results

1. Held-out main sample: 0/10 complete matched A/B/C blocks; 0/30 saved main contexts (30 not received), and 0/30 complete 24-week games. Main inference is withheld until the fixed sample and provenance gates pass.
2. Development is separate: 4 latency attempts and 12 feasibility attempts. Development choices are assessed by completion, timing, display adherence and knowledge; they do not enter main hypothesis tests.
3. A context controls one four-role chain. There are ten planned independent matched contrasts, not forty independent students. These are model observations under fictional classroom instructions; human participants: 0.
4. Stopped prefixes, zero choices, opposite responses and failures are retained. Partial cost is not full-horizon cost; missing decisions are never filled.

Requested model/effort: gpt-6-astra / xhigh. Actually reported labels appear per case in the outcome table; unverified labels remain unverified. Launcher temperature and reasoning-token allocation are not exposed.

Figures and status | Case outcomes | Validation and exact input hashes

Main hypotheses and quantities

Endpoint	Quantity and units	Contrast; predicted sign
H1	How much the new order falls for each extra case already on the way; average of eight paired scenario responses, in ordered cases per outstanding case	B-A; negative
H2	Total holding + backlog cost over four roles and 24 weeks; game cost units per chain	B-A; positive
H3	The same response to outstanding orders, comparing educated and control students while both face the shorter deadline	C-B; positive
H4a	How much each role's orders fluctuate across 24 weeks: sample order variance, averaged over four roles; cases squared	C-B; negative
H4b	Order variance divided by external customer-demand variance, averaged over roles: the bullwhip ratio; dimensionless	C-B; negative

Holding costs 0.5 per remaining case per week; backlog costs 1 per unfilled case per week. H4b is undefined for zero demand variance. H4a/H4b share a numerator and are not independent replications. Role outcomes include final observed inventory, backlog and outstanding after the last accepted order; consequences after week 24 are unmeasured.

Test	Available pairs	Mean contrast	Pointwise 95% interval	Two-sided p	Holm, family of 5	Direction/status
H1	0	NA	NA to NA	NA	NA	unavailable; inference withheld
H2	0	NA	NA to NA	NA	NA	unavailable; inference withheld
H3	0	NA	NA to NA	NA	NA	unavailable; inference withheld
H4a	0	NA	NA to NA	NA	NA	unavailable; inference withheld
H4b	0	NA	NA to NA	NA	NA	unavailable; inference withheld

NA means unavailable or deliberately withheld. Before the whole fixed sample and external reviews pass, any available-pair effects are descriptive only; treatment-related noncompletion can bias completer comparisons. There is no completer-only causal claim and no significance-based stopping or replacement.

Current gate conditions: Frozen MAIN_PLAN.json with exactly ten unique block_ids is missing or invalid; Independent fresh-context block assignment is not confirmed in the frozen plan; Main plan is not marked frozen; The selected V1/V2 wording is not frozen in the main plan; All ten assigned A/B/C blocks have not completed replay-valid 24-week games and all probes; Thirty passing external adherence reviews bound to current RESULT and CASE hashes are required.

When eligible, tests reuse the separately frozen paired sign-flip implementation; the entire five-endpoint family receives Holm correction. A significant opposite effect is reported as opposite evidence. Intervals are unadjusted paired-block percentile bootstrap intervals (10,000 resamples), not simultaneous intervals. The sign-flip null requires independent symmetric/exchangeable paired differences; a hash receipt alone cannot verify independence or this assumption. Each block contrast.

Completion, timing and knowledge checks

Stage	Saved attempts	Complete cases	Complete 24-week games	Replay-valid	Passing external reviews
latency	4	4	0	4	4
feasibility	12	5	0	12	12
main	0	0	0	0	0

Development qualification checks (fractions range from 0 to 1):

Stage/version/arm	Complete cases	On time / timed displays	Game median (s)	Probe median (s)	Knowledge fraction
feasibility/V1/A	2/2	1	3.50375	3.44525	1
feasibility/V1/B	1/2	0.947368	3.48825	3.54275	1
feasibility/V1/C	0/2	0.857143	3.5505	3.4845	NA
feasibility/V2/A	2/2	1	3.5665	3.6015	1
feasibility/V2/B	0/2	0.8	3.461	3.4605	NA
feasibility/V2/C	0/2	0.8	3.33175	NA	NA
latency/V1/A	2/2	NA	3.598	3.5155	1
latency/V2/A	2/2	NA	4.0935	3.5745	1

Prompt qualification by version/stage/arm reports complete fractions, on-time accepted responses per opened timed display, game/probe case-median and maximum latencies, and post-task knowledge. Untimed calibration has no on-time fraction. Latencies measure display-open to complete-payload reception, including transport, scheduling and output; they are not thinking time. A stopped case can have zero late submission events when no payload arrives. Open decisions remain pending until the controller persists their expiry.

All submission timings retain rejections. Post-behavior item scores use private answer keys; scores never exclude main cases. The report does not automatically declare a development wording qualified or select one.

Probe interpretation

Pair slopes, role summaries, and pair-fixed-effect fits/residuals preserve all observed responses. The pipeline aggregate weights eight pairs equally. The backlog comparison uses (q_highB - q_lowB)/(B_high - B_low) over eight separate pairs, also in ordered cases per backlog case. A missing pair makes the fixed aggregate unavailable. Fixed-effect regression is supplementary and weights pair slopes by squared quantity differences; it does not replace the primary equal-pair mean.

Report zero-choice and zero-slope fractions, order minima/maxima, role-level slopes and cross-pair spread without deleting extremes. Two levels per pair cannot identify curvature; cross-regime variation is a descriptive heterogeneity diagnostic. The pipeline slope is a responsiveness proxy, not an identified normative underweighting coefficient. Flat choices across both dimensions cannot establish pipeline-specific neglect. Known assay arrival schedules differ from uncertain gameplay fulfillment.

Planning and reproducibility

Minimum attainable p versus independent chain blocks is a mathematical planning figure, not observed evidence or power. With ten nonzero pairs the exact two-sided minimum is 2/1024 = 0.001953125; zero differences/ties may raise it. The fixed N=10 sample does not promise significance.

Rebuild offline: python scripts/report_exp14.py. It scans only data/{latency,feasibility,main}/*/RESULT.json, never other experiments, never an API, and never a hidden reasoning trace. It checks controller replay, exact CASE/result binding and external review receipts. A main external ADHERENCE_REVIEW.json must contain passed: true, result_sha256 (exact file bytes), and case_sha256 (controller canonical digest). Its full_public_context_reviewed, fresh_context_confirmed, literal_model_choices_confirmed, and fixed_requested_settings_confirmed checks must all be true. hypotheses/MAIN_PLAN.json must freeze exactly ten block_ids, the selected prompt_version, and independent_blocks_confirmed: true with frozen: true. Reviews must assess full-display access, literal choices, isolation and source/reading exposure without using effect direction or p-values.

Current input fingerprint: 0f45aa5491fbb00237e9af93b5584f2519467ec6e4ae77278105a36078809054. Changed-input regenerations first preserve prior generated bytes under reports/history/<fingerprint>/ and

analysis/history/<fingerprint>/. Current generated-file manifest. The local event hash chain is reproducibility evidence, not tamper-proof authentication.