Experiment 14: current results

1. Held-out main sample: 0/10 complete matched A/B/C blocks; 0/30 saved main contexts (30 not received), and 0/30 complete 24-week games. Main inference is withheld until the fixed sample and provenance gates pass.
2. Development is separate: 3 latency attempts and 0 feasibility attempts. Development choices are assessed by completion, timing, display adherence and knowledge; they do not enter main hypothesis tests.
3. A context controls one four-role chain. There are ten planned independent matched contrasts, not forty independent students. These are model observations under fictional classroom instructions; human participants: 0.
4. Stopped prefixes, zero choices, opposite responses and failures are retained. Partial cost is not full-horizon cost; missing decisions are never filled.

Requested model/effort: gpt-6-astra / xhigh. Actually reported labels appear per case in the outcome table; unverified labels remain unverified. Launcher temperature and reasoning-token allocation are not exposed.

Figures and status | Case outcomes | Validation and exact input hashes

Main hypotheses and quantities

Endpoint	Quantity and units	Contrast; predicted sign
H1	Mean eight pipeline-pair slopes (q_lowP - q_highP)/(P_high - P_low); ordered cases per outstanding case	B-A; negative
H2	Total holding + backlog cost over four roles and 24 weeks; game cost units per chain	B-A; positive
H3	The same pipeline response proxy as H1	C-B; positive
H4a	Mean of four role-level sample weekly order variances (denominator 23); cases squared	C-B; negative
H4b	Mean role order variance divided by the chain's external-customer sample demand variance; dimensionless	C-B; negative

Holding costs 0.5 per remaining case per week; backlog costs 1 per unfilled case per week. H4b is undefined for zero demand variance. H4a/H4b share a numerator and are not independent replications. Role outcomes include final observed inventory, backlog and outstanding after the last accepted order; consequences after week 24 are unmeasured.

Test	Available pairs	Mean contrast	Pointwise 95% interval	Two-sided p	Holm, family of 5	Direction/status
H1	0	NA	NA to NA	NA	NA	unavailable; inference withheld
H2	0	NA	NA to NA	NA	NA	unavailable; inference withheld
H3	0	NA	NA to NA	NA	NA	unavailable; inference withheld
H4a	0	NA	NA to NA	NA	NA	unavailable; inference withheld
H4b	0	NA	NA to NA	NA	NA	unavailable; inference withheld

NA means unavailable or deliberately withheld. Before the whole fixed sample and external reviews pass, any available-pair effects are descriptive only; treatment-related noncompletion can bias completer comparisons. There is no completer-only causal claim and no significance-based stopping or replacement.

Current gate conditions: Frozen MAIN_PLAN.json with exactly ten unique block_ids is missing or invalid; Independent fresh-context block assignment is not confirmed in the frozen plan; Main plan is not marked frozen; The selected V1/V2 wording is not frozen in the main plan; All ten assigned A/B/C blocks have not completed replay-valid 24-week games and all probes; Thirty passing external adherence reviews bound to current RESULT and CASE hashes are required.

When eligible, tests reuse the separately frozen paired sign-flip implementation; the entire five-endpoint family receives Holm correction. A significant opposite effect is reported as opposite evidence. Intervals are unadjusted paired-block percentile bootstrap intervals (10,000 resamples), not simultaneous intervals. The sign-flip null requires independent symmetric/exchangeable paired differences; a hash receipt alone cannot verify independence or this assumption. Each block contrast.

Completion, timing and knowledge checks

| Stage | Saved attempts | Complete cases | Complete 24-week games | Replay-valid | Passing external reviews |

latency	3	2	0	3	0
feasibility	0	0	0	0	0
main	0	0	0	0	0

Development qualification checks (fractions range from 0 to 1):

Stage/version/arm	Complete cases	On time / timed displays	Game median (s)	Probe median (s)	Knowledge fraction
latency/V1/A	1/2	NA	3.18	3.633	1
latency/V2/A	1/1	NA	3.6405	3.469	1

Prompt qualification by version/stage/arm reports complete fractions, on-time accepted responses per opened timed display, game/probe case-median and maximum latencies, and post-task knowledge. Untimed calibration has no on-time fraction. Latencies measure display-open to complete-payload reception, including transport, scheduling and output; they are not thinking time. A stopped case can have zero late submission events when no payload arrives. Open decisions remain pending until the controller persists their expiry.

All submission timings retain rejections. Post-behavior item scores use private answer keys; scores never exclude main cases. The report does not automatically declare a development wording qualified or select one.

Probe interpretation

Pair slopes, role summaries, and pair-fixed-effect fits/residuals preserve all observed responses. The pipeline aggregate weights eight pairs equally. The backlog comparison uses (q_highB - q_lowB)/(B_high - B_low) over eight separate pairs, also in ordered cases per backlog case. A missing pair makes the fixed aggregate unavailable. Fixed-effect regression is supplementary and weights pair slopes by squared quantity differences; it does not replace the primary equal-pair mean.

Report zero-choice and zero-slope fractions, order minima/maxima, role-level slopes and cross-pair spread without deleting extremes. Two levels per pair cannot identify curvature; cross-regime variation is a descriptive heterogeneity diagnostic. The pipeline slope is a responsiveness proxy, not an identified normative underweighting coefficient. Flat choices across both dimensions cannot establish pipeline-specific neglect. Known assay arrival schedules differ from uncertain gameplay fulfillment.

Planning and reproducibility

Minimum attainable p versus independent chain blocks is a mathematical planning figure, not observed evidence or power. With ten nonzero pairs the exact two-sided minimum is 2/1024 = 0.001953125; zero differences/ties may raise it. The fixed N=10 sample does not promise significance.

Rebuild offline: python scripts/report_exp14.py. It scans only data/{latency,feasibility,main}/*/RESULT.json, never other experiments, never an API, and never a hidden reasoning trace. It checks controller replay, exact CASE/result binding and external review receipts. A main external ADHERENCE_REVIEW.json must contain passed: true, result_sha256 (exact file bytes), and case_sha256 (controller canonical digest). Its full_public_context_reviewed, fresh_context_confirmed, literal_model_choices_confirmed, and fixed_requested_settings_confirmed checks must all be true. hypotheses/MAIN_PLAN.json must freeze exactly ten block_ids, the selected prompt_version, and independent_blocks_confirmed: true with frozen: true. Reviews must assess full-display access, literal choices, isolation and source/reading exposure without using effect direction or p-values.

Current input fingerprint: 3918baaf5d2dcf4d5274f450ad769eff8bcea88079dbda882266f74c1ecd7197. Changed-input regenerations first preserve prior generated bytes under reports/history/<fingerprint>/ and analysis/history/<fingerprint>/. Current generated-file manifest. The local event hash chain is reproducibility evidence, not tamper-proof authentication.