Simulated student time pressure

1. **The treatment is fictional student time-pressure framing: 120 seconds in A versus 8 seconds in B/C.** The model has no actual runtime or response-admission deadline. Every scored game order and probe answer is an accepted model choice; no defaults or carry-forward actions are used.

2. Development: 0/3 fresh contexts saved, 0/3 complete 24-week games, 0/1 complete matched block. Development effects are descriptive; all development p-values are NA.

3. Main: 0/30 contexts saved; **0/10 complete matched blocks**. Main hypothesis inference remains withheld until the full fixed sample and external reviews pass.

4. This measures model responses to a classroom framing, not real student behavior, actual thinking time, or a verified cognitive time limit. Four roles and 24 weeks are repeated observations within a context. The independent main sample is ten matched chain blocks, not forty students.

A receives the control reading and the 120-second student framing. B receives the same control reading and the 8-second framing. C receives the education reading and the same 8-second framing as B. All arms use one accepted choice per decision; post-task questions follow the game and 32 probes. Requested model/effort is gpt-6-astra / xhigh. Actual reported labels remain separate; temperature and hidden reasoning-token allocation are not exposed.

Figures | Case outcomes | Validation

Five hypothesis quantities

Endpoint	Plain definition and units	Contrast; predicted sign
H1	How much the order falls per extra case already on the way, averaged equally over eight fixed pipeline pairs; ordered cases per outstanding case	B-A; negative
H2	Holding plus backlog cost across four roles and all 24 weeks; game cost units per chain	B-A; positive
H3	The same pipeline-response measure, comparing education with control under the 8-second student framing	C-B; positive
H4a	Sample variance of weekly orders within each role (denominator 23), averaged over four roles; order cases squared	C-B; negative
H4b	Role order variance divided by the chain's external-customer sample demand variance, averaged over roles; dimensionless bullwhip ratio	C-B; negative

Holding costs 0.5 per remaining case per week; backlog costs 1 per unfilled case per week. Partial trajectories have no complete-horizon H2/H4 estimate. H4b is undefined if customer-demand variance is zero. H4a/H4b share their numerator and are not independent confirmations. The pipeline assay is a responsiveness proxy, not an identified normative underweighting coefficient; flat answers in both pipeline and backlog assays do not establish pipeline-specific neglect.

Development contrasts

Test	Pairs	Signed contrast	95% interval	Raw two-sided p	Holm 5	Direction
H1 B-A	0	NA	NA to NA	NA	NA	NA
H2 B-A	0	NA	NA to NA	NA	NA	NA
H3 C-B	0	NA	NA to NA	NA	NA	NA
H4a C-B	0	NA	NA to NA	NA	NA	NA
H4b C-B	0	NA	NA to NA	NA	NA	NA

Main contrasts

Test	Pairs	Signed contrast	95% interval	Raw two-sided p	Holm 5	Direction
H1 B-A	0	NA	NA to NA	NA	NA	NA
H2 B-A	0	NA	NA to NA	NA	NA	NA
H3 C-B	0	NA	NA to NA	NA	NA	NA
H4a C-B	0	NA	NA to NA	NA	NA	NA
H4b C-B	0	NA	NA to NA	NA	NA	NA

NA means no complete paired endpoint or inference deliberately withheld. Available-pair contrasts before the full main gate are descriptive only; they do not justify a completer-only causal claim. Significant opposite effects, if observed in the completed main sample, remain opposite evidence.

Completion, comprehension and source checks

Stage / arm / block	Weeks	Pipeline pairs / 8	Post-task score	Replay	Review
No saved framing cases	NA	NA	NA	NA	NA

All readings, hypothetical states, actual choices and public transcripts remain source records. Pair/role diagnostics retain zero choices, zero slopes, extremes and separate backlog response. Pair slopes keep all eight fixed pipeline pairs; an absent pair makes the aggregate unavailable rather than changing its weights. Two levels within a pair cannot identify curvature. Comprehension is scored after behavior; low scores never select individual main cases. Recorded wall latency includes transport and scheduling and is observational, never a response-admission cutoff.

Fixed main inference and prefix curves

Only the complete frozen ten-block main sample receives primary two-sided paired sign-flip inference, with Holm correction across five endpoints. Pointwise 95% intervals use 10,000 paired-block bootstrap resamples and are not simultaneous intervals. Independent symmetric/exchangeable paired differences are required by the sign-flip null. A receipt cannot establish that assumption or authenticate hidden reasoning.

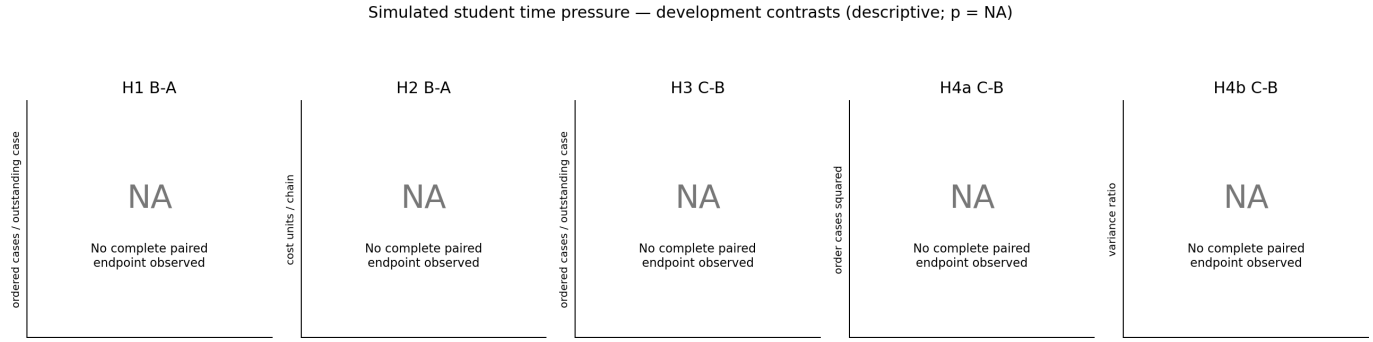
Chain-prefix curves use the frozen order at n=1 through 10; n=10 is primary. Prefix effects use only fully available prefixes, and raw prefix p-values appear only after the entire main gate passes. n<2 has p=NA. These are descriptive sensitivity plots, not predicted significance crossings, power estimates, or a rule to extend, stop, replace, or select a sample. No favorable outcome, slope, direction or p-value is used for wording selection or main-case inclusion.

Reproduction

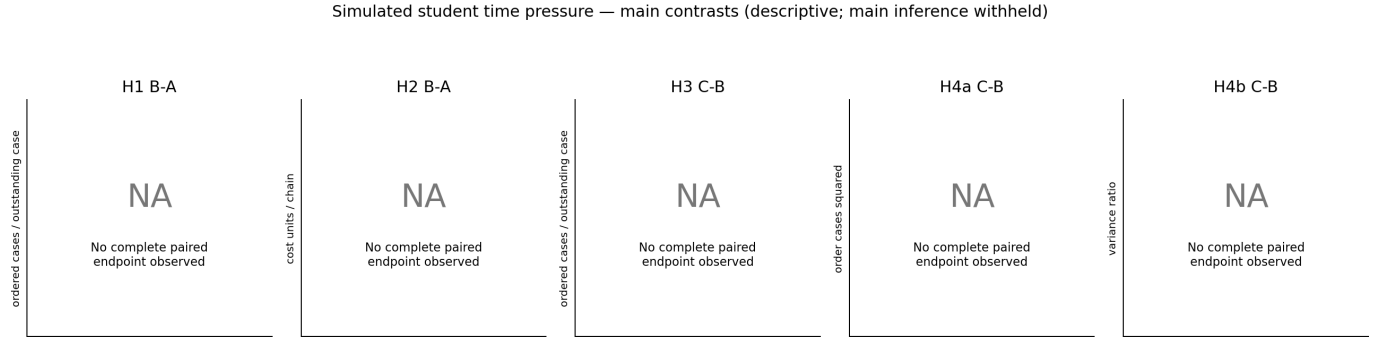
Rebuild: .venv/Scripts/python.exe scripts/report_exp14_framing.py. Only data/student_framing_v3/development and data/student_framing_v3/main are read. The strict-deadline study and unused carry-forward draft stay separate. No API requests, participant launches, fabricated missing decisions, or hidden reasoning traces are used. Main requires STUDENT_FRAMING_V3_MAIN_PLAN.json, thirty complete replay-valid actual-response contexts and hash-bound external reviews. Full gate reasons and input hashes.

Input fingerprint: 844966806147c0bad21f48cd949999bc9628893b252681b8b78a289f61601818. Earlier generated bytes are preserved in this amendment's history folders before regeneration.

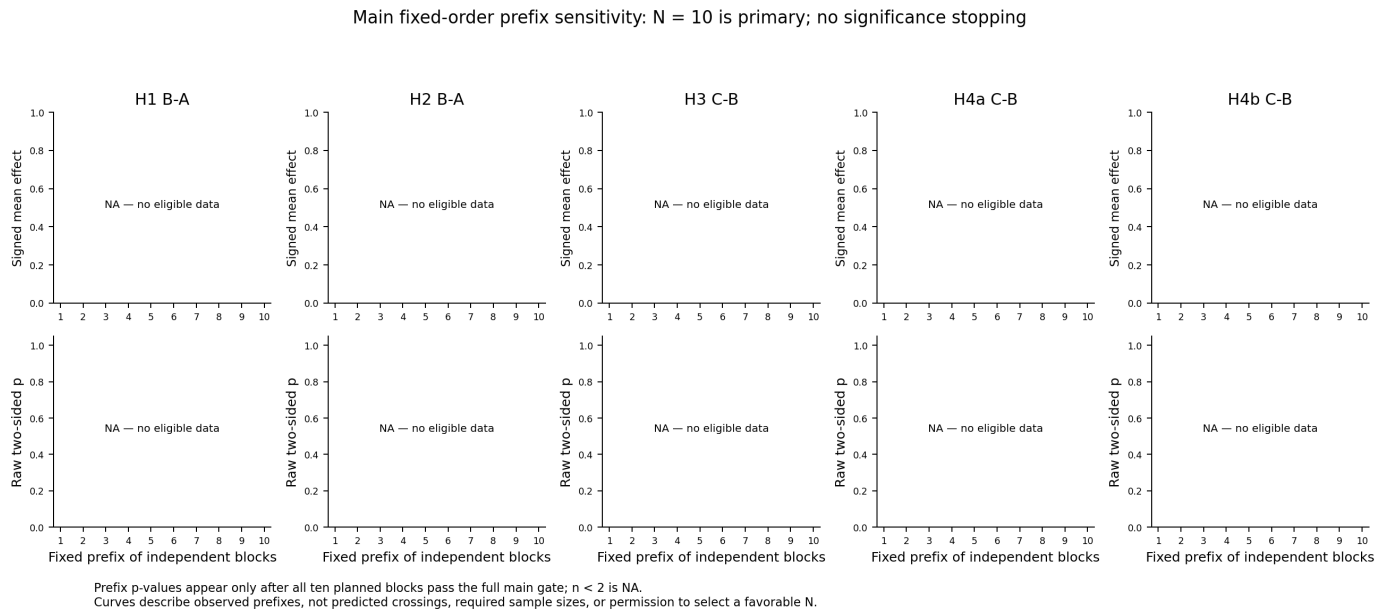
Development figures



Each point is one matched block; NA contains no fabricated observation. All development p-values remain NA.



The full ten-block sample is required for primary main inference.



The frozen n=10 sample is primary. Curves do not predict a significant crossing or authorize a stopping change.