Experiment 14: H1-H4 conclusions

GPT-6 Astra / extra high. Ten matched chain pairs; 40 fictional roles per condition; 24 weeks. Development responses are excluded.

1. H1: Pressure - adequate, pipeline response: mean difference -0.1228; Holm p=0.009766.

2. H2: Pressure - adequate, total holding/backlog cost: mean difference +68.5; Holm p=0.009766.

3. H3: Education - control, pipeline response: mean difference +0.07411; Holm p=0.009766.

4. H4: Education - control: order variance -7.725 (Holm p=0.009766); bullwhip -1.117 (Holm p=0.009766).

H	Quantity / units	Mean	95% interval	Raw p	Holm 5
H1	Pipeline response (Order cases / outstanding case)	-0.1228	-0.1384 to -0.1076	0.001953	0.009766
H2	Total cost (Game cost units)	+68.5	45.5 to 96.55	0.001953	0.009766
H3	Pipeline response (Order cases / outstanding case)	+0.07411	0.0516 to 0.09498	0.001953	0.009766
H4a	Order variance (Cases squared)	-7.725	-11.72 to -4.333	0.001953	0.009766
H4b	Bullwhip (Ratio)	-1.117	-1.666 to -0.6332	0.001953	0.009766

H1/H2 compare pressure with adequate time. H3/H4 compare education with control under pressure. Tests are exact two-sided paired sign flips; Holm adjusts the five endpoints. Intervals are pointwise 95% paired-block bootstrap intervals. All ten paired differences have the predicted sign for every endpoint, so each exact raw p equals 2/1024; the identical p-values are expected.

These are AI student-roleplay results, not observations of real students. Actual model runtime was unlimited; every answer came from the model. Pipeline response is a proxy for underweighting. Shared project context and the block6 service interruption are documented in the full report/context note.